

# Objective

## Multi-threading

- To understand how multiple threads can execute in parallel
- To learn how to implement threads
- To understand race conditions and deadlocks
- To be able to avoid corruption of shared objects by using locks and conditions
- To be able to use threads for programming animations

Study Chapter 22 (Web Only) : Early Object 7<sup>th</sup> editon.

# Threads

- A thread is a program that is executed independently of other parts of the program
- The Java Virtual Machine executes each thread in the program for a short amount of time (Time Slice)
- This gives the impression of parallel execution

# Creating and Running a Thread

Develop a class that implements the Runnable interface

```
public interface Runnable
{
    void run();
}
```

Place the code for your task into the run method of your class

```
public class MyRunnable implements Runnable
{
    public void run()
    {
        // Task statements go here
        . . .
    }
}
```

# Running a Thread

Create an object of your subclass

```
Runnable r = new MyRunnable();
```

Construct a `Thread` object from the runnable object.

```
Thread t = new Thread(r);
```

Call the `start` method to start the thread.

```
t.start();
```

# GreetingRunnable Outline

```
public class GreetingRunnable implements Runnable
{
    public GreetingRunnable(String aGreeting)
    {
        greeting = aGreeting;
    }

    public void run()
    {
        // Task statements go here
        . . .
    }
    // Fields used by the task statements
    private String greeting;
}
```

# Example

## Example 1

Thu	Dec	28	23:12:03	PST	2004	Hello,	World!
Thu	Dec	28	23:12:04	PST	2004	Hello,	World!
Thu	Dec	28	23:12:05	PST	2004	Hello,	World!
Thu	Dec	28	23:12:06	PST	2004	Hello,	World!
Thu	Dec	28	23:12:07	PST	2004	Hello,	World!
Thu	Dec	28	23:12:08	PST	2004	Hello,	World!
Thu	Dec	28	23:12:09	PST	2004	Hello,	World!
Thu	Dec	28	23:12:10	PST	2004	Hello,	World!
Thu	Dec	28	23:12:11	PST	2004	Hello,	World!
Thu	Dec	28	23:12:12	PST	2004	Hello,	World!

Output

# Thread Actions

To wait a thread, use the sleep method of the Thread class

```
sleep(milliseconds)
```

A sleeping thread can generate an **InterruptedException**

- Catch the exception
- Terminate the thread

# Generic run Method

```
public void run()
{
    try
    {
        Task statements
    }
    catch (InterruptedException exception)
    {
    }
    finally{
        //Clean up, if necessary
    }
}
```



# File GreetingRunnable.java

```
01: import java.util.Date;
02:
03: /**
04:     A runnable that repeatedly prints a greeting.
05: */
06: public class GreetingRunnable implements Runnable
07: {
08:     /**
09:         Constructs the runnable object.
10:         @param aGreeting the greeting to display
11:     */
12:     public GreetingRunnable(String aGreeting)
13:     {
14:         greeting = aGreeting;
15:     }
16:
17:     public void run()
18:     {
```

# File GreetingRunnable.java

```
19:         try
20:         {
21:             for (int i = 1; i <= REPETITIONS; i++)
22:             {
23:                 Date now = new Date();
24:                 System.out.println(now + " " + greeting);
25:                 Thread.sleep(DELAY);
26:             }
27:         }
28:         catch (InterruptedException exception)
29:         {
30:         }
31:     }
32:
33:     private String greeting;
34:
35:     private static final int REPETITIONS = 10;
36:     private static final int DELAY = 1000;
37: }
```

# To Start the Thread

- Construct an object of your runnable class

```
Runnable r = new GreetingRunnable("Hello World");
```

- Then construct a thread and call the `start` method.

```
Thread t = new Thread(r);  
t.start();
```

# Extend Thread

Another method to create and run a thread:

```
public class MyThread extends Thread{  
    public void run(){  
        // Task statements go here  
    }  
}
```

Then construct an object of the class and call start() method.

Check example 2

```
Thread t = new MyThread();  
t.start();
```

# File GreetingThreadTester.java

```
01: import java.util.Date;
02:
03: /**
04:     This program tests the greeting thread by running two
05:     threads in parallel.
06: */
07: public class GreetingThreadTester
08: {
09:     public static void main(String[] args)
10:     {
11:         Thread t1 = new GreetingRunnable("Hello, World!");
12:         Thread t2 = new GreetingRunnable("Goodbye, World!");
13:         t1.start();
14:         t2.start();
15:
16:     }
17: }
18: }
```

**What would be the output of this program?**

# Thread Scheduler

- The thread scheduler runs each thread for a short amount of time (a *time slice*)
- Then the scheduler activates another thread
- There will always be slight variations in running times especially when calling operating system services (e.g. input and output)
- There is no guarantee about the order in which threads are executed

# Self Check

- What happens if you change the call to the `sleep` method in the `run` method to `Thread.sleep(1)`?
- What would be the result of the program if the `main` method called

```
r1.run();  
r2.run();
```

instead of starting threads?

# Answers

- The messages are printed about one millisecond apart.
- The first call to `run` would print 10 "Hello" messages, and then the second call to `run` would print 10 "Goodbye" messages



# Pool of Threads

If you have to create many short-lived threads, then use pool of threads.

```
Runnable r1 = new MyThread(...);  
Runnable r2 = new MyThread(...);  
ExecutorService pool = Executors.newFixedThreadPool(MAX);  
pool.execute(r1);  
pool.execute(r2);
```

This is an efficient method of creating threads for a server side program.

# Terminating Threads

- A thread terminates when its **run** method terminates
- Do not terminate a thread using the deprecated **stop** method. Why?
- Instead, notify a thread that it should terminate

```
t.interrupt();
```

- **interrupt** does not cause the thread to terminate—it sets a boolean field in the thread data structure

# Terminating Threads

The run method should check occasionally if it has been interrupted

- Use the `interrupted` method
- An interrupted thread should release resources, clean up, and exit

```
public void run()
{
    for (int i = 1;
        i <= REPETITIONS && !Thread.interrupted(); i++)
    {
        // Do work
    }
    // Clean up
}
```

# Terminating Threads

The sleep method throws an InterruptedException when a sleeping thread is interrupted

- Catch the exception
- Terminate the thread

```
public void run() {  
    try {  
        for (int i = 1; i <= REPETITIONS; i++) {  
            //Do work  
        }  
    }  
    catch (InterruptedException exception) {  
    }  
    finally{  
        // Clean up  
    }  
}
```

# Terminating Threads

- Java does not force a thread to terminate when it is interrupted
- It is entirely up to the thread what it does when it is interrupted
- Interrupting is a general mechanism for getting the thread's attention

# Self Check

Consider the following runnable.

```
public class MyRunnable implements Runnable {  
    public void run() {  
        try {  
            System.out.println(1);  
            Thread.sleep(1000);  
            System.out.println(2);  
        }  
        catch (InterruptedException exception) {  
            System.out.println(3);  
        }  
        System.out.println(4);  
    }  
}
```

# Self Check

Suppose a thread with this runnable is started and immediately interrupted.

```
Thread t = new Thread(new MyRunnable());  
t.start();  
t.interrupt();
```

What output is produced?

## Answers

The `run` method prints the values 1, 3, and 4.

The call to `interrupt` merely sets the interruption flag, but the `sleep` method immediately throws an

**`InterruptedException`**.

Note that this interrupt does not terminate the program.



# Race Conditions

- When threads share a common object, they can conflict with each other
- Sample program: multiple threads manipulate a bank account
- Create a BankAccount object
- Create two threads:
  - t1 deposits \$100 into the bank account for 10 iterations
  - t2 withdraws \$100 from the bank account for 10 iterations

# Race Conditions

Here is the run method of `DepositRunnable`:

```
public void run(){
    try {
        for (int i = 1; i <= count; i++) {
            account.deposit(amount);
            Thread.sleep(DELAY);
        }
    }
    catch (InterruptedException exception) {
    }
}
```

The `WithdrawRunnable` class is similar

# Sample Application

The result should be something like this:

```
Depositing 100.0, new balance is 100.0
Withdrawing 100.0, new balance is 0.0
Depositing 100.0, new balance is 100.0
Depositing 100.0, new balance is 200.0
Withdrawing 100.0, new balance is 100.0
. . .
Withdrawing 100.0, new balance is 0.0
```

Example 3

- But sometimes you may find that balance is not correct.
- like this:

Why?

```
Depositing 100.0Withdrawing 100.0,
    new balance is 100.0, new balance is -100.0
```

# Analysis

Thread: 1

```
deposit(double M){  
    newB = b + M;  
    b = newB;  
}
```

Thread: 2

```
withdraw(double M){  
    newB = b - M;  
    b = newB;  
}
```

Assume balance (b) of bank account is 100, and  $M = 100$

T1 runs: newB = 200 and then process changes T2 by OS.

T2 runs: b still is 100, and so newB =  $100 - 100 = 0$ , and then process changes changes to T1 by OS.

T1 runs: b = newB → b = 200; process changes changes to T2 by OS.

T2 runs: b = newB → b = 0;

# Analysis

Assume we add a flag to bank account, and process check this flag if it is being used by other process.

Thread: 1

```
deposit(double M){  
    while(flag!=0)  
        continue;  
    flag = 1;  
    newB = b + M;  
    b= newB;  
    flag = 0;  
}
```

Thread: 2

```
withdraw(double M){  
    while(flag!=0)  
        continue;  
    flag = 1;  
    newB = b - M;  
    b = newB;  
    flag = 0;  
}
```

Do you think this will work?

# Race Condition

- Occurs if the effect of multiple threads on shared data depends on the order in which they are scheduled
- It is possible for a thread to reach the end of its time slice in the middle of a statement
- It may evaluate the right-hand side of an equation but not be able to store the result until its next turn

# Synchronizing Object Access

- To solve this problem such as the one just seen, use a *lock object*
- A lock object is used to control threads that manipulate shared resources
- In Java: **Lock** interface and several classes that implement it
  - **ReentrantLock**: most commonly used lock class
  - Locks are a feature of Java version 5.0
  - Earlier versions of Java have a lower-level facility for thread **synchronization**

# Synchronizing Object Access

- Typically, a lock object is added to a class whose methods access shared resources, like this:

```
public class BankAccount
{
    public BankAccount()
    {
        balanceChangeLock = new ReentrantLock();
        . . .
    }
    . . .
    private Lock balanceChangeLock;
}
```



# Synchronizing Object Access

- Code that manipulates shared resource is surrounded by calls to **lock** and **unlock**:

```
balanceChangeLock.lock();  
//Code that manipulates the shared resource  
balanceChangeLock.unlock();
```

# Synchronizing Object Access

- If code between calls to **lock** and **unlock** throws an exception, call to **unlock** **never happens**
- To overcome this problem, place call to **unlock** into a **finally** clause:

# Synchronizing Object Access

```
public void deposit(double amount)
{
    balanceChangeLock.lock();
    try
    {
        System.out.print("Depositing " + amount);
        double newBalance = balance + amount;
        System.out.println(", new balance is " + newBalance);
        balance = newBalance;
    }
    finally
    {
        balanceChangeLock.unlock();
    }
}
```

Example 4

# Synchronizing Object Access

- When a thread calls `lock`, it owns the lock until it calls `unlock`
- A thread that calls `lock` while another thread owns the lock is temporarily deactivated
- Thread scheduler periodically reactivates thread so it can try to acquire the lock
- Eventually, waiting thread can acquire the lock

# Self Check

---

- If you construct two `BankAccount` objects, how many lock objects are created?
- What happens if we omit the call `unlock` at the end of the `deposit` method?

# Answers

---

- Two, one for each bank account object. Each lock protects a separate balance field.
- When a thread calls `deposit`, it continues to own the lock, and any other thread trying to deposit or withdraw money in the same bank account is blocked forever.

# Deadlocks

T1:

```
deposit(double M){  
    bLock.lock();  
    try    {  
        b=b+M;  
    }  
    finally {  
        bLock.unlock();  
    }  
}
```

Assume initial  
balance (b) is 0

T2:

```
withdraw((double M){  
    bLock.lock();  
    try    {  
        while (b<=0)  
            continue;  
        b=b-M;  
    }  
    finally {  
        bLock.unlock();  
    }  
}
```

1. T2 runs, and lock the bLock, but b is zero and it stuck in the loop.
2. T1 runs, and it blocks because T2 locked in.
3. Both threads are **deadlocked**.

# Avoiding Deadlocks

---

- A deadlock occurs if no thread can proceed because each thread is waiting for another to do some work first



# Condition Objects

- To overcome problem, use a **condition object**
- Condition objects allow a thread to temporarily release a lock, and to regain the lock at a later time
- Each condition object belongs to a specific lock object

# Condition Objects

You obtain a condition object with `newCondition` method of `Lock` interface

```
public class BankAccount
{
    public BankAccount()
    {
        balanceChangeLock = new ReentrantLock();
        sufficientFundsCondition
            = balanceChangeLock.newCondition();
        . . .
    }
    . . .
    private Lock balanceChangeLock;
    private Condition sufficientFundsCondition;
}
```

# Condition Objects

- It is customary to give the condition object a name that describes condition to test
- You need to implement an appropriate test
- As long as test is not fulfilled, call `await` on the condition object:

# Condition Objects

```
public void withdraw(double amount)
{
    balanceChangeLock.lock();
    try
    {
        while (balance < amount)
            sufficientFundsCondition.await();
        . . .
    }
    finally
    {
        balanceChangeLock.unlock();
    }
}
```

# Condition Objects

- Calling `await`
  - Makes current thread wait
  - Allows another thread to acquire the lock object
- To unlock, another thread must execute `signalAll` *on the same condition object*

```
sufficientFundsCondition.signalAll();
```

- `signalAll` unblocks all threads waiting on the condition

Example 5

# Condition Objects

- `signalAll`: randomly picks just one thread waiting on the object and unblocks it
- `signal` can be more efficient, but you need to know that every waiting thread can proceed
- Recommendation: always call `signalAll`

# Self Check

What is the essential difference between calling `sleep` and `await`?

Why is the `sufficientFundsCondition` object a field of the `BankAccount` class and not a local variable of the `withdraw` and `deposit` methods?

# Answers

A sleeping thread is reactivated when the sleep delay has passed. A waiting thread is only reactivated if another thread has called `signalAll` or `signal`.

The calls to `await` and `signal/signalAll` must be made *to the same object*.



# Animation

Animation can be done much better using threads.  
Each object runs independently.

Example 6

# Summary

A thread is a program unit that is executed concurrently with other parts of the program.

The start method of the Thread class starts a new thread that executes the run method of the associated Runnable object.

The sleep method puts the current thread to sleep for a given number of milliseconds.

When a thread is interrupted, the most common response is to terminate the run method.

The thread scheduler runs each thread for a short amount of time, called a time slice.

A thread terminates when its run method terminates.

# Summary

The run method can check whether its thread has been interrupted by calling the interrupted method.

A race condition occurs if the effect of multiple threads on shared data depends on the order in which the threads are scheduled.

By calling the lock method, a thread acquires a Lock object. Then no other thread can acquire the lock until the first thread releases the lock.

A deadlock occurs if no thread can proceed because each thread is waiting for another to do some work first.

# Summary

Calling `await` on a condition object makes the current thread wait and allows another thread to acquire the lock object.

A waiting thread is blocked until another thread calls `signalAll` or `signal` on the condition object for which the thread is waiting.